

Exclusive Group Lasso

Coleman Zhang

May 10, 2019

1 Introduction

Assume design matrix $\mathbf{X} \in \mathbb{R}^{n \times p}$, response $\mathbf{y} \in \mathbb{R}^n$. Let us consider penalized linear regression, where we are minimizing the objective:

$$\operatorname{argmin}_{\beta} \frac{1}{2} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \Omega(\beta) \quad (1)$$

where $\lambda \in \mathbb{R}_{++}$, and $\Omega(\cdot) : \mathbb{R}^p \rightarrow \mathbb{R}_+$ defined as

$$\Omega(\beta) = \frac{1}{2} \sum_{g \in \mathcal{G}} \omega_g \|\beta_g\|_1^2 = \frac{1}{2} \sum_{g \in \mathcal{G}} \omega_g \left(\sum_{j \in g} |\beta_j| \right)^2 \quad (2)$$

is the exclusive group lasso norm. The group weights should be a positive vector, i.e.,

$$\omega_g \geq 0 \quad \forall g \in \mathcal{G}$$

A convention introduced in the group lasso is to pick the group weight to be proportional to the group size: e.g. $\omega_g := n_g$. In addition, we limit group assignments to be mutually exclusive, and collectively exhaustive, i.e.,

$$g_k \cap g_l = \emptyset \quad \forall k \neq l, \quad \bigcup_k g_k = \{1, \dots, p\} = [p]$$

In this case one could prove that $\Omega(\cdot)$ is a norm. In addition, our objective is convex, but not differentiable at $\beta_i = 0$ for some i .

Definition 1. (*Vector Norm*) Let $\|\cdot\| : \mathbb{C}^m \rightarrow \mathbb{R}_+$. Then $\|\cdot\|$ is norm if $\forall x, y \in \mathbb{C}^m$ and $\forall \alpha \in \mathbb{C}$

- $x \neq 0 \implies \|x\| > 0$ (*positive definite*);
- $\|\alpha x\| = |\alpha| \|x\|$ (*homogeneous*);
- $\|x + y\| \leq \|x\| + \|y\|$ (*triangle inequality*).

2 Literature

This work relies on extensive results from convex analysis and optimization theory. For a review, see Boyd & Vandenberghe [BV04, V20], Beck [B17], and Tibshirani's lecture notes [T18]. In particular, the ability to

derive the dual program using the KKT conditions, taking the Fenchel conjugate of functions, and computing the dual norm is especially important.

Kong et. al [KF14, KL16] provided an iterative re-weighted least square algorithm. Sun et. al [SC20] built on this work and provided a bisection algorithm that solves the lasso problem at each iteration. Campbell & Allen [CA15] provided a coordinate descent method and extensive theoretical analysis including a result for the dual norm of $\Omega(\cdot)$.

However, none of these work provides an algorithm that adopts both the state-of-the-art coordinate descent, and convergence checks based on strong duality. Xiang et. al [XW14] and Fercoq et. al [FG15], for example, provided screening rules to safely set coefficients to 0 using duality of the lasso. Ndiaye et. al [NF16, NF17, WY14] generalizes the "duality gap safe" screening rule first to the sparse group lasso [SF13], and finally to a general sparsity-enforcing penalty. We follow the methodologies of the previously mentioned work for the following reasons. First, (block) coordinate descent achieves the state-of-the-art performance in terms of speed of convergence. Second, we could check if the algorithm converges and early stop by computing the duality gap.

3 Primal, Dual, & Optimality

Let $P_\lambda(\beta)$ be the primal objective,

$$\hat{\beta}^{(\lambda)} \in \operatorname{argmin}_{\beta \in \mathbb{R}^p} P_\lambda(\beta) = \operatorname{argmin}_{\beta \in \mathbb{R}^p} \frac{1}{2} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \Omega(\beta)$$

This is equivalent to,

$$\begin{aligned} \text{(P)} \quad & \min_{\beta \in \mathbb{R}^p, \mathbf{z} \in \mathbb{R}^n} \frac{1}{2} \|\mathbf{y} - \mathbf{z}\|_2^2 + \frac{\lambda}{2} \sum_{g \in \mathcal{G}} \omega_g \|\beta_g\|_1^2 \\ & \text{s.t. } \mathbf{z} = \mathbf{X}\beta \end{aligned} \tag{3}$$

The Lagrangian dual of this problem is,

$$\begin{aligned} & \min_{\beta \in \mathbb{R}^p, \mathbf{z} \in \mathbb{R}^n} \frac{1}{2} \|\mathbf{y} - \mathbf{z}\|_2^2 + \frac{\lambda}{2} \sum_{g \in \mathcal{G}} \omega_g \|\beta_g\|_1^2 + \mathbf{u}^T (\mathbf{z} - \mathbf{X}\beta) \\ & = \min_{\mathbf{z} \in \mathbb{R}^n} \frac{1}{2} \|\mathbf{y} - \mathbf{z}\|_2^2 + \mathbf{u}^T \mathbf{z} + \min_{\beta \in \mathbb{R}^p} \lambda \sum_{g \in \mathcal{G}} \frac{1}{2} \omega_g \|\beta_g\|_1^2 - \mathbf{u}^T \mathbf{X}\beta \\ & = \frac{1}{2} \|\mathbf{y}\|_2^2 - \frac{1}{2} \|\mathbf{y} - \mathbf{u}\|_2^2 - \lambda \max_{\beta \in \mathbb{R}^p} \left(\frac{\mathbf{X}^T \mathbf{u}}{\lambda} \right)^T \beta - \Omega(\beta) \\ & = \frac{1}{2} \|\mathbf{y}\|_2^2 - \frac{1}{2} \|\mathbf{y} - \mathbf{u}\|_2^2 - \lambda \Omega^* \left(\frac{\mathbf{X}^T \mathbf{u}}{\lambda} \right) \end{aligned} \tag{4}$$

where Ω^* is the Fenchel conjugate of Ω ,

Definition 2. (Conjugate functions) Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$. The function $f^* : \mathbb{R}^n \rightarrow \mathbb{R}$ defined as

$$f^*(\mathbf{y}) = \sup_{x \in \operatorname{dom} f} (\mathbf{y}^T \mathbf{x} - f(\mathbf{x}))$$

is called the conjugate of function f .

Property. The conjugate function has the following calculus rules:

- (Scaling) Let $f : \mathbb{E} \rightarrow (-\infty, \infty]$ and let $\alpha \in \mathbb{R}_{++}$. The conjugate of $g(\mathbf{x}) = \alpha f(\mathbf{x})$ is given by

$$g^*(\mathbf{y}) = \alpha f^*\left(\frac{\mathbf{y}}{\alpha}\right)$$

- (Separable functions) Let $g : \mathbb{E}_1 \times \mathbb{E}_2 \times \cdots \times \mathbb{E}_p \rightarrow (-\infty, \infty]$ be given by $g(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_p) = \sum_{i=1}^p f_i(\mathbf{x}_i)$, where $f_i : \mathbb{E}_i \rightarrow (-\infty, \infty]$ is a proper function for any $i = 1, 2, \dots, p$. Then

$$g^*(\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_p) = \sum_{i=1}^p f_i^*(\mathbf{y}_i)$$

Definition 3. (Dual norm) Let $\|\cdot\|$ be a norm on \mathbb{R}^n . The associated dual norm, denoted $\|\cdot\|_*$, is defined as

$$\|\mathbf{z}\|_* = \sup_{\|\mathbf{x}\| \leq 1} \mathbf{z}^T \mathbf{x}$$

From the definition we have the following Hölder's inequality,

$$\mathbf{z}^T \mathbf{x} \leq \|\mathbf{x}\| \|\mathbf{z}\|_*$$

Example. (Norm squared). Now consider the function $f(\mathbf{x}) = \frac{1}{2} \|\mathbf{x}\|^2$, where $\|\cdot\|$ is a norm, with dual norm $\|\cdot\|_*$. From Hölder's inequality,

$$\mathbf{y}^T \mathbf{x} - \frac{1}{2} \|\mathbf{x}\|^2 \leq \|\mathbf{x}\| \|\mathbf{y}\|_* - \frac{1}{2} \|\mathbf{x}\|^2$$

for all \mathbf{x} . The righthand side is a quadratic function of $\|\mathbf{x}\|$, which has maximum $\frac{1}{2} \|\mathbf{y}\|_*^2$. Therefore for all \mathbf{x} , we have

$$\mathbf{y}^T \mathbf{x} - \frac{1}{2} \|\mathbf{x}\|^2 \leq \frac{1}{2} \|\mathbf{y}\|_*^2$$

which shows that $f^*(\mathbf{y}) \leq \frac{1}{2} \|\mathbf{y}\|_*^2$.

To show the other inequality, pick \mathbf{x} to be such that $\mathbf{y}^T \mathbf{x} = \|\mathbf{y}\|_* \|\mathbf{x}\|$, scaled so that $\|\mathbf{x}\| = \|\mathbf{y}\|_*$. Then, from definition and by construction,

$$f^*(\mathbf{y}) \geq \mathbf{y}^T \mathbf{x} - \frac{1}{2} \|\mathbf{x}\|^2 = \frac{1}{2} \|\mathbf{y}\|_*^2.$$

In conclusion, it must be that $f^*(\mathbf{y}) = \frac{1}{2} \|\mathbf{y}\|_*^2$

Proposition 1. Let

$$\Omega(\boldsymbol{\beta}) = \frac{1}{2} \sum_{g \in \mathcal{G}} \omega_g \|\boldsymbol{\beta}_g\|_1^2 = \frac{1}{2} \sum_{g \in \mathcal{G}} \omega_g \left(\sum_{j \in g} |\beta_j| \right)^2$$

Then the conjugate of Ω is

$$\Omega^*(\boldsymbol{\xi}) = \frac{1}{2} \sum_{g \in \mathcal{G}} \frac{1}{\omega_g} \|\boldsymbol{\xi}_g\|_\infty^2 = \frac{1}{2} \sum_{g \in \mathcal{G}} \frac{1}{\omega_g} \left(\max_{j \in g} |\xi_j| \right)^2$$

Proof. First by observation $\mathbf{\Omega}$ is group separable. Define $\mathbf{\Omega}(\boldsymbol{\beta}) = \sum_{g \in \mathcal{G}} \mathbf{\Omega}_g(\boldsymbol{\beta}_g) = \sum_{g \in \mathcal{G}} \frac{1}{2} \omega_g \|\boldsymbol{\beta}_g\|_1^2$, then $\mathbf{\Omega}^*(\boldsymbol{\xi}) = \sum_{g \in \mathcal{G}} \mathbf{\Omega}_g^*(\boldsymbol{\xi}_g)$. Define $f(\boldsymbol{\beta}_g) = \frac{1}{2} \|\boldsymbol{\beta}_g\|_1^2$, then $\mathbf{\Omega}_g(\boldsymbol{\beta}_g) = \omega_g f(\boldsymbol{\beta}_g)$. By the previous example and the calculus rule for scaling, we have $\mathbf{\Omega}_g^*(\boldsymbol{\xi}_g) = \frac{1}{2} \omega_g \|\frac{\boldsymbol{\xi}_g}{\omega_g}\|_\infty^2 = \frac{1}{2} \frac{1}{\omega_g} \|\boldsymbol{\xi}_g\|_\infty^2$. Hence $\mathbf{\Omega}^*(\boldsymbol{\xi}) = \sum_{g \in \mathcal{G}} \mathbf{\Omega}_g^*(\boldsymbol{\xi}_g) = \frac{1}{2} \sum_{g \in \mathcal{G}} \frac{1}{\omega_g} \|\boldsymbol{\xi}_g\|_\infty^2$. \square

From the above illustration, the dual formulation is given by,

$$(D) \max_{\mathbf{u} \in \mathbb{R}^n} \frac{1}{2} \|\mathbf{y}\|_2^2 - \frac{1}{2} \|\mathbf{y} - \mathbf{u}\|_2^2 - \frac{1}{2\lambda} \sum_{g \in \mathcal{G}} \frac{1}{\omega_g} \|\mathbf{X}_g^T \mathbf{u}\|_\infty^2 =: D_\lambda(\mathbf{u}) \quad (5)$$

By Slater's condition, strong duality applies. If $\boldsymbol{\beta}^*$ and \mathbf{u}^* are both primal and dual optimal, then there is no duality gap, i.e., $P_\lambda(\boldsymbol{\beta}^*) = D_\lambda(\mathbf{u}^*)$.

4 Subgradient, KKT conditions

We establish the following optimality condition based on subgradient:

Proposition 2. (Fermat's Rule) (see (Bauschke and Combettes, 2011, Proposition 26.1) for a more general result) For any convex function $f: \mathbb{R}^d \rightarrow \mathbb{R}$,

$$x^* \in \underset{x}{\operatorname{argmin}} f(x) \implies 0 \in \partial f(x^*)$$

From the sufficient conditions for optimality, if $\boldsymbol{\beta}^*$ and \mathbf{u}^* satisfy KKT conditions (Fermat's Rule + primal dual feasible), then they are primal and dual optimal, which implies zero duality gap. In the following section, we derive the subgradient equations for the exclusive lasso problem.

Recall,

$$(P) \min_{\boldsymbol{\beta}} \frac{1}{2} \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 + \frac{\lambda}{2} \sum_{g \in \mathcal{G}} \omega_g \left(\sum_{j \in g} |\beta_j| \right)^2$$

The subgradient equation for a particular β_i satisfies:

$$\mathbf{X}_i^T (\mathbf{r}_{-i} - \mathbf{X}_i \beta_i) = \lambda \omega_g \left(\sum_{k \in g \setminus i} |\beta_k| + |\beta_i| \right) \partial |\beta_i| \quad (6)$$

where

$$\mathbf{r}_{-i} = \mathbf{y} - \sum_{j \neq i} \mathbf{X}_j \beta_j$$

is the partial residual, and subgradient of ℓ_1 norm

$$\partial |\beta_i| = \begin{cases} -1 & \text{if } \beta_i < 0 \\ [-1, 1] & \text{if } \beta_i = 0 \\ 1 & \text{if } \beta_i > 0 \end{cases} \quad (7)$$

For the lasso problem, the subgradient equations are crucial for deriving the soft-threshold operator used in coordinate-wise descent. For a brief review, please see Tibshirani's lecture notes. Here it is possible to derive a closed form solution. Define the soft-threshold operator (at level $\tau \geq 0$) for $\mathbf{x} \in \mathbb{R}^d$ as $[\mathcal{S}_\tau(\mathbf{x})]_j = \text{sign}(x_j)(|x_j| - \tau)_+$,

$$\begin{aligned}\tilde{\tau} &= \frac{\lambda\omega_g \sum_{k \in g \setminus i} |\beta_k|}{\lambda\omega_g + \|\mathbf{X}_i\|_2^2} \\ \tilde{z} &= \frac{\mathbf{X}_i^T \mathbf{r}_{-i}}{\lambda\omega_g + \|\mathbf{X}_i\|_2^2} \\ \hat{\beta}_i &= \mathcal{S}_{\tilde{\tau}}(\tilde{z})\end{aligned}\tag{8}$$

5 Stopping criterion

We stop according the following criterion,

$$P_\lambda(\hat{\beta}) - D_\lambda(\mathbf{r}) < \epsilon$$

which reflects the tolerance on the duality gap.

6 Bonus: Sparse Exclusive Group Lasso

Consider the penalty

$$\begin{aligned}\Omega(\beta) &= \Omega_1^{1-\alpha}(\beta) + \Omega_2^\alpha(\beta) \\ &= \frac{1}{2}(1-\alpha) \sum_{g \in \mathcal{G}} \omega_g \|\beta_g\|_1^2 + \alpha \|\beta\|_1\end{aligned}\tag{9}$$

where $\alpha \in [0, 1]$.

Theorem 1. (conjugate of sum). Let $h_1 : \mathbb{E} \rightarrow (-\infty, \infty]$ be a proper convex function and $h_2 : \mathbb{E} \rightarrow \mathbb{R}$ be a real-valued convex function. Then

$$(h_1 + h_2)^* = h_1^* \square h_2^*$$

where \square denotes the infimal convolution operator, defined as

$$(f_1 \square f_2)(\mathbf{x}) = \inf_{\mathbf{u} \in \mathbb{E}} f_1(\mathbf{u}) + f_2(\mathbf{x} - \mathbf{u})$$

We know that,

$$\begin{aligned}(\Omega_1^{1-\alpha})^*(\xi) &= \frac{1}{2} \sum_{g \in \mathcal{G}} \frac{1}{(1-\alpha)\omega_g} \|\xi_g\|_\infty^2 \\ (\Omega_2^\alpha)^*(\xi) &= \mathbf{I}_{\mathcal{B}_\infty}(\xi/\alpha) = \sum_{g \in \mathcal{G}} \mathbf{I}_{\mathcal{B}_\infty}(\xi_g/\alpha)\end{aligned}\tag{10}$$

By definition we have,

$$\begin{aligned}
\Omega^*(\xi) &= ((\Omega_1^{1-\alpha})^* \square (\Omega_2^\alpha)^*)(\xi) \\
&= \inf_{\eta} (\Omega_1^{1-\alpha})^*(\xi - \eta) + (\Omega_2^\alpha)^*(\eta) \\
&= \sum_{g \in \mathcal{G}} \inf_{\eta_g} \frac{1}{2} \frac{1}{(1-\alpha)\omega_g} \|\xi_g - \eta_g\|_\infty^2 + \mathbf{I}_{\mathcal{B}_\infty}(\eta_g/\alpha) \\
&= \sum_{g \in \mathcal{G}} \inf_{\|\eta_g\|_\infty \leq \alpha} \frac{1}{2} \frac{1}{(1-\alpha)\omega_g} \|\xi_g - \eta_g\|_\infty^2
\end{aligned} \tag{11}$$

which is equivalent to solving

$$\begin{aligned}
\nu_g^* &= \min_{\eta_g} \|\xi_g - \eta_g\|_\infty \\
&\text{s.t. } \|\eta_g\|_\infty \leq \alpha
\end{aligned} \tag{12}$$

We can see that $\eta_g^*(\xi_g)$ is indeed the projection of ξ on $\alpha\mathcal{B}_\infty$, which admits closed form solution:

$$[\eta_g^*(\xi_g)]_i = [\mathbf{P}_{\alpha\mathcal{B}_\infty}(\xi_g)]_i = \begin{cases} \alpha & \text{if } [\xi_g]_i > \alpha \\ [\xi_g]_i & \text{if } |[\xi_g]_i| \leq \alpha \\ -\alpha & \text{if } [\xi_g]_i < -\alpha \end{cases} \tag{13}$$

Hence (12) can be solved as

$$\begin{aligned}
\nu_g^* &= \|\mathcal{S}_\alpha(\xi_g)\|_\infty \\
\Omega^*(\xi) &= \sum_{g \in \mathcal{G}} \frac{1}{2} \frac{1}{(1-\alpha)\omega_g} \|\mathcal{S}_\alpha(\xi_g)\|_\infty^2
\end{aligned} \tag{14}$$

The soft-threshold operator for the sparse exclusive group lasso problem becomes:

$$\begin{aligned}
\tilde{\tau} &= \frac{\lambda(1-\alpha)\omega_g \sum_{k \in g \setminus i} |\beta_k| + \lambda\alpha}{\lambda(1-\alpha)\omega_g + \|\mathbf{X}_i\|_2^2} \\
\tilde{z} &= \frac{\mathbf{X}_i^T \mathbf{r}_{-i}}{\lambda(1-\alpha)\omega_g + \|\mathbf{X}_i\|_2^2} \\
\hat{\beta}_i &= \mathcal{S}_{\tilde{\tau}}(\tilde{z})
\end{aligned} \tag{15}$$

Design choice: not choosing

$$\|\beta\|_{1,2} = \sqrt{\sum_{g \in \mathcal{G}} \omega_g \|\beta_g\|_1^2}$$

with dual norm [B17]

$$\|\beta\|_{1,2}^D = \sqrt{\sum_{g \in \mathcal{G}} \frac{1}{\omega_g} \|\beta_g\|_\infty^2}$$

since the penalty is not coordinate-wise separable, so it might be problematic to apply coordinate descent algorithm.

Question: What is λ_{\max} ? How to determine the dual scaling constant to find the dual feasible point?

6.1 Group Level & Feature Level Screening

For λ large enough, $\mathbf{0} \in \partial P_\lambda(\boldsymbol{\beta})$. Using first-order conditions, we can determine λ_{\max} . For a particular lambda, we can determine when a particular $\beta_i = 0$ or $\boldsymbol{\beta}_g = \mathbf{0}$.

Feature level screening:

$$\mathbf{X}_j^T (\mathbf{r}_{-j} - \mathbf{X}_j \boldsymbol{\beta}_j) = \lambda \alpha \partial |\beta_j| + \lambda (1 - \alpha) \omega_g \|\boldsymbol{\beta}_g\|_1 \partial |\beta_j|$$

$$\forall j \in g, |\mathbf{X}_j^T \mathbf{r}_{-j}| < \lambda \left(\alpha + (1 - \alpha) \omega_g \sum_{k \in g \setminus j} |\beta_k| \right) \implies \hat{\beta}_j = 0.$$

Group level screening:

$$\mathbf{X}_g^T (\mathbf{r}_{-g} - \mathbf{X}_g \boldsymbol{\beta}_g) = \lambda \alpha \partial \|\boldsymbol{\beta}_g\|_1 + \lambda (1 - \alpha) \omega_g \|\boldsymbol{\beta}_g\|_1 \partial \|\boldsymbol{\beta}_g\|_1$$

$$\forall g \in \mathcal{G}, \|\mathbf{X}_g^T \mathbf{r}_{-g}\|_\infty < \lambda \alpha \implies \hat{\boldsymbol{\beta}}_g = \mathbf{0}.$$

$$\lambda_{\max} = \max_{g \in \mathcal{G}} \|\mathbf{X}_g^T \mathbf{r}_{-g}\|_\infty / \alpha$$

$$\forall \lambda > \lambda_{\max}, \hat{\boldsymbol{\beta}} = \mathbf{0}$$

References

- [BV04] S. BOYD and L. VANDENBERGHE “Convex Optimization,” Cambridge University Press, 2004.
- [B17] A. BECK “First-Order Methods in Optimization,” Society for Industrial and Applied Mathematics, 2017.
- [T18] R. TIBSHIRANI “Convex Optimization Fall 2018,” <https://www.stat.cmu.edu/~ryantibs/convexopt-F18>, 2018.
- [V20] L. VANDENBERGHE “ECE236C - Optimization Methods for Large-Scale Systems,” <http://www.seas.ucla.edu/~vandenbe/ee236c.html>, 2020.
- [KF14] D. KONG and R. FUJIMAKI and J. LIU and F. NIE and C. CHRIS “Exclusive Feature Learning on Arbitrary Structures via $\ell_{1,2}$ -norm,” Advances in Neural Information Processing Systems 27, 2014.
- [KL16] D. KONG and J. LIU and B. LIU and X. BAO “Uncorrelated Group LASSO,” AAAI, 2016.
- [SC20] Y. SUN and B. CHAIN and S. KASKI and J. SHAWE-TAYLOR “Correlated Feature Selection with Extended Exclusive Group Lasso,” 2020.
- [CA15] F. CAMPBELL and G. ALLEN “Within Group Variable Selection through the Exclusive Lasso,” Electronic Journal of Statistics, 2015.
- [FG15] O. FERCOQ and A. GRAMFORT and J. SALMON “Mind the duality gap: safer rules for the Lasso,” Proceedings of Machine Learning Research, 2015.
- [NF16] E. NDIAYE and O. FERCOQ and A. GRAMFORT and J. SALMON “GAP Safe Screening Rules for Sparse-Group Lasso,” Advances in Neural Information Processing Systems 29, 2016.
- [NF17] E. NDIAYE and O. FERCOQ and A. GRAMFORT and J. SALMON “Gap Safe Screening Rules for Sparsity Enforcing Penalties,” Journal of Machine Learning Research, 2017.
- [WY14] J. WANG and J. YE “Two-Layer Feature Reduction for Sparse-Group Lasso via Decomposition of Convex Sets,” Journal of Machine Learning Research, 2014.
- [XW14] Z. XIANG and Y. WANG and P. RAMADGE “Screening Tests for Lasso Problems,” IEEE Transactions on Pattern Analysis and Machine Intelligence, 2014.
- [SF13] N. SIMON and J. FRIEDMAN and T. HASTIE and R. TIBSHIRANI “A sparse-group lasso,” Journal of Computational and Graphical Statistics, 2013.